


class_07 Machine Learning

Iulia Rusu

Main function for kmeans clustering in base R is called 'kmeans()'

```
hist(rnorm(5000, mean= 3))
```

Histogram of rnorm(5000, mean = 3)




```
# Make a vector with 60 total points half centered at +3 and half centered at -3
tmp <- c(rnorm(30, mean =3), rnorm(30, mean= -3))
#reversed a vector to generate another vector with temp
tmp1 <- rev(tmp)
#cbind to bind to vectors as columns, generates matrix
temperature <- cbind(tmp, tmp1)
temperature
```

	tmp	tmp1
[1,]	3.579611	-3.964540
[2,]	2.140252	-2.966405
[3,]	4.389818	-2.059953
[4,]	3.509826	-4.132507
[5,]	2.703410	-3.116287
[6,]	1.366379	-3.757341
[7,]	1.763849	-2.294419
[8,]	4.277852	-3.762414
[9,]	2.391975	-4.160339
[10,]	3.633930	-3.302875
[11,]	1.337579	-1.906407
[12,]	2.958263	-3.133425
[13,]	2.156720	-4.153809
[14,]	4.209870	-3.587902
[15,]	1.607313	-3.908292
[16,]	3.895723	-2.664850
[17,]	4.699163	-3.278570
[18,]	2.880225	-3.722995
[19,]	3.919782	-4.219083
[20,]	3.762063	-3.250798
[21,]	1.988400	-3.634352
[22,]	3.288278	-3.035885
[23,]	1.976397	-5.371975
[24,]	3.392614	-3.033089
[25,]	4.571973	-3.557470
[26,]	4.978979	-4.317606
[27,]	3.685530	-3.878485
[28,]	0.944076	-2.716172
[29,]	3.304277	-1.927053
[30,]	2.698474	-2.811844
[31,]	-2.811844	2.698474
[32,]	-1.927053	3.304277
[33,]	-2.716172	0.944076
[34,]	-3.878485	3.685530
[35,]	-4.317606	4.978979
[36,]	-3.557470	4.571973
[37,]	-3.033089	3.392614
[38,]	-5.371975	1.976397
[39,]	-3.035885	3.288278
[40,]	-3.634352	1.988400
[41,]	-3.250798	3.762063
[42,]	-4.219083	3.919782

```
[43,] -3.722995  2.880225  
[44,] -3.278570  4.699163  
[45,] -2.664850  3.895723  
[46,] -3.908292  1.607313  
[47,] -3.587902  4.209870  
[48,] -4.153809  2.156720  
[49,] -3.133425  2.958263  
[50,] -1.906407  1.337579  
[51,] -3.302875  3.633930  
[52,] -4.160339  2.391975  
[53,] -3.762414  4.277852  
[54,] -2.294419  1.763849  
[55,] -3.757341  1.366379  
[56,] -3.116287  2.703410  
[57,] -4.132507  3.509826  
[58,] -2.059953  4.389818  
[59,] -2.966405  2.140252  
[60,] -3.964540  3.579611
```

```
plot(temperature)
```



```
k <- kmeans(temperature, centers=2, nstart = 20)
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	tmp	tmp1
1	3.067087	-3.387572
2	-3.387572	3.067087

Clustering vector:

[illegible]

Within cluster sum of squares by cluster:

```
[1] 53.19723 53.19723
      (between_SS / total_SS =  92.2 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

```
attributes(k)
```

\$names

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

```
$class
```

```
[1] "kmeans"
```

```
#What are the cluster centers
k$centers
```

k\$centers


	tmp	tmp1
1	3.067087	-3.387572
2	-3.387572	3.067087

k\$cluster

[illegible]

Plot data as 'x' showing your clustering result and the center point for each cluster?


```
plot( temperature, col=k$cluster)
points(k$centers, pch=15, col="green")
```



```
#recycle property, allows you to make new data
```

```
#Run kmeans and cluster into 3 groups, plot
```

```
k3 <- kmeans(temperature, centers = 3, nstart = 20)
plot(temperature, col=k3$cluster)
```



```
k$tot.withinss
```

```
[1] 106.3945
```

```
k3$tot.withinss
```

```
[1] 80.04027
```

Big limitation of kmeans is that it imposes structure on your data that you ask for in the first place.

#Heirarchical Clustering

The main function in “base” R for this is called ‘hclust()’ It wants a distance matrix as input not the data itself

We can calculate a distance matrix in lots of different ways bit here we will use the ‘dist()’ function.

```
d <- dist(temperature)
hc <- hclust(d)
```

```
hc
```

Call:


```
hclust(d = d)
```

Cluster method : complete

Distance : euclidean

Number of objects: 60

```
plot(hc)
abline(h=9, col="red")
```




There is a special plot method for hclust objects. Let's see it.

To get the cluster membership vector we need to “cut” the tree at a given height that we [ick]

```
grps <- cutree(hc, k=2)
grps
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

```
plot(temperature, col=grps)
```




Principle Component Analysis

```
url <- "https://tinyurl.com/UK-foods"
x1 <- read.csv(url, row.names = 1)
```

```
head(x1)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139


```
#One useful plot in this case (beause we only have 4 countries to look accross) is a "pairs
pairs(x1, col=rainbow(10), pch=16)
```



##Enter PCA

The main function to do a PCA in “base” R is called ‘prcomp’ It wants our foods as the columns and the countries as the rows. It basically wants the transpose of the table.

```
#transpose with t(x)

pca <- prcomp(t(x1))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

```
attributes(pca)
```

```
$names
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

```
$class
[1] "prcomp"
```

```
pca$x1
```

```
NULL
```

```
pca$rotation
```

	PC1	PC2	PC3	PC4
Cheese	-0.056955380	0.016012850	0.02394295	-0.409382587
Carcass_meat	0.047927628	0.013915823	0.06367111	0.729481922
Other_meat	-0.258916658	-0.015331138	-0.55384854	0.331001134
Fish	-0.084414983	-0.050754947	0.03906481	0.022375878
Fats_and_oils	-0.005193623	-0.095388656	-0.12522257	0.034512161
Sugars	-0.037620983	-0.043021699	-0.03605745	0.024943337
Fresh_potatoes	0.401402060	-0.715017078	-0.20668248	0.021396007
Fresh_Veg	-0.151849942	-0.144900268	0.21382237	0.001606882
Other_Veg	-0.243593729	-0.225450923	-0.05332841	0.031153231
Processed_potatoes	-0.026886233	0.042850761	-0.07364902	-0.017379680
Processed_Veg	-0.036488269	-0.045451802	0.05289191	0.021250980
Fresh_fruit	-0.632640898	-0.177740743	0.40012865	0.227657348
Cereals	-0.047702858	-0.212599678	-0.35884921	0.100043319
Beverages	-0.026187756	-0.030560542	-0.04135860	-0.018382072
Soft_drinks	0.232244140	0.555124311	-0.16942648	0.222319484
Alcoholic_drinks	-0.463968168	0.113536523	-0.49858320	-0.273126013
Confectionery	-0.029650201	0.005949921	-0.05232164	0.001890737

```
plot(pca$x[, 1], pca$x[,2], xlab = "PC1 (67%)", ylab="PC2 (29%)",
     col=c("orange", "red", "blue", "darkgreen"), pch=15)
abline(h=0, col="gray", lty=2)
abline(v=0, col="gray", lty=2)
```

